

SRE Practices: Alerting and Observability Framework

1. Segregation of Alerts Based on Severity

To ensure effective incident management and reduce alert fatigue, alerts will be categorized by severity:

- **Severity 0 (Sev 0): Critical Alerts**
 - Description: Issues requiring immediate resolution to avoid significant business impact.
 - Examples:
 - Breached critical thresholds (e.g., CPU utilization > 90%).
 - Certificate expirations.
 - Service outages or unavailability.
 - Action: Immediate escalation to on-call teams via designated distribution lists.
- **Severity 1 (Sev 1): High-Priority Alerts**
 - Description: High-priority issues needing prompt attention but not immediately business-critical.
 - Examples:
 - Moderate performance degradation (e.g., API latency > 500ms consistently).
 - Resource nearing critical thresholds (e.g., disk space > 80%).
 - Action: Escalation to on-call teams with clear timelines for resolution.
- **Severity 2 (Sev 2): Informational Alerts**
 - Description: Alerts triggered by events within specific time frames or patterns for monitoring purposes.
 - Examples:
 - High volume of requests within a short window.
 - Spikes in non-critical errors (e.g., 404s or 429s).
 - Action: Monitored by teams; no immediate escalation unless thresholds persist.

2. Classification of Alerts: Actionable vs. Non-Actionable

- **Actionable Alerts**
 - Require immediate investigation and corrective actions.
 - Examples:
 - Pods unreachable or unhealthy.
 - Database connection failures.

- **Non-Actionable Alerts**

- For monitoring or informational purposes only; no action required.
- Examples:
 - Successful job completions.
 - Scheduled report generation.
- Action: Limit distribution to essential stakeholders to minimize noise.

3. Standardized Alert Pattern

A consistent naming convention will improve clarity and facilitate alert triaging. The proposed format is:

<Severity>-<Application Name>-<Description>

- Example: Sev0-OrderService-HighCPUUsage

4. Observability Framework with New Relic and Splunk

To achieve end-to-end observability, we will leverage New Relic and Splunk across the three pillars of observability:

- **Metrics**

- **New Relic:** Monitor application performance metrics, infrastructure health, and key business transactions.
 - Dashboards: CPU usage, memory utilization, request latency.
 - Alerts: Threshold-based alerts for anomalies (e.g., high error rates).
- **Splunk:** Aggregate and visualize system-level metrics for historical analysis.

- **Logs**

- **New Relic:** Centralized log management for real-time application logging.
 - Use cases: Debugging and tracing application errors.
- **Splunk:** Advanced log search and analysis.
 - Use cases: Root cause analysis, compliance, and historical trend identification.

- **Traces**

- **New Relic:** Distributed tracing for end-to-end visibility of application workflows.
 - Use cases: Identifying bottlenecks and slow API calls.
- **Splunk:** Augmented with tracing data to correlate logs and metrics for comprehensive insights.

5. Supporting Documentation

- **New Relic:**

- Guidelines for configuring alerts and dashboards.
 - Best practices for anomaly detection and APM setup.
- **Splunk:**
 - Documentation for building custom dashboards and alert rules.
 - Recommended search queries for efficient log analysis.